

# Sponsors for Global Azure Bootcamp

Thanks !



# Azure HDInsight

- Azure Bootcamp -

---

2017. 4. 22

Ingyu Lee  
(inlee@troy.edu)



<https://github.com/Azure-Bootcamp-Troy/AzureHDInsight>

# Contents

---

- Target Audience
  - ✓ Database / BI Professional.
  - ✓ Data Scientists / Analysts with some technical experience.
- Prerequisites
  - ✓ Familiarity with database concepts and basic SQL query syntax.
  - ✓ Basic understanding of Unix commands.
  - ✓ Familiarity with programming fundamentals.
- Outline
  - P1: Getting Started with [HDInsight](#).
  - P2: Processing Big Data with [Hive](#).
  - P3: Going Beyond Hive with [Pig](#) and Python.
  - P4: Using [HBase](#) for NoSQL Data
- Hands-on Labs
  - ✓ Free trial available.
  - ✓ Microsoft Azure Subscription.
- edX (Microsoft)
  - ✓ [DAT202.1x](#) Processing Big Data with Hadoop in Azure HDInsight
  - ✓ [DAT202.2x](#) Implementing Real-Time Analytics with Hadoop in Azure HDInsight
  - ✓ [DAT202.3x](#) Implementing Predictive Analytics with Spark in Azure HDInsight
- MVA (Microsoft Virtual Academy)
  - ✓ Big Data Analytics with HDInsight
  - ✓ Implementing Big Data Analysis Jump Start
- Coursera (Big Data Specialization)
  - ✓ Hadoop Platform and Application Framework
  - ✓ Introduction to Big Data Analytics

# Free eBooks

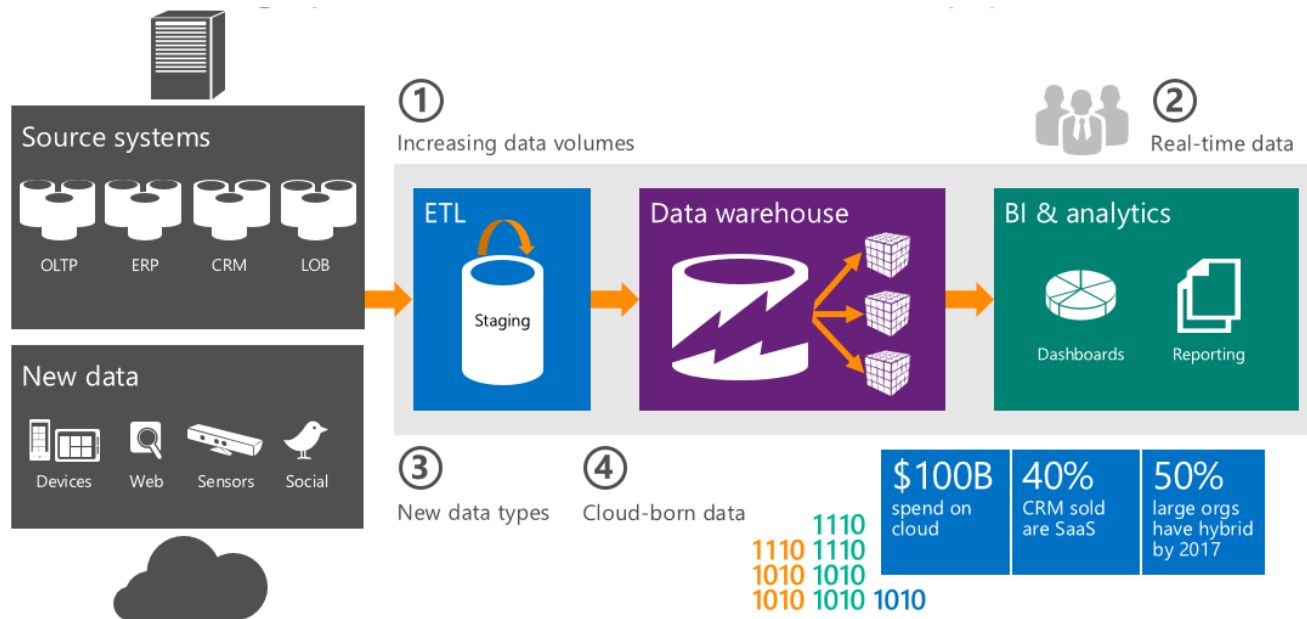


<http://www.microsoftvirtualacademy.com/ebooks>

# I Introduction to Big Data and HDInsight

## 1. What is Big Data?

- a. Data that is **too large or complex for analysis** in traditional relational databases.
- b. Typified by the “3 Vs”
  - a. **Volume**: Huge amounts of data to process.
  - b. **Variety**: A mixture of structured and unstructured data.
  - c. **Velocity**: New data generated extremely frequently.




**Breaking Point of Traditional Approach**

# I Introduction to Big Data and HDInsight: What is Hadoop?


## 1. What is Hadoop?

- Open source distributed data processing cluster.
- Data processed in Hadoop Distributed File System (HDFS).
- Resource Management is performed by YARN.




Apache Open Source Project  
Highly scalable distributed file system (HDFS)  
Distributed processing on data nodes


Data volumes




Data variety



Data velocity





Hadoop is a platform with portfolio of projects

Governed by Apache Software Foundation (ASF)


Comprises core services of MapReduce, HDFS, and YARN


In addition to the core, includes functions across:

Data services which allow you to manipulate and move data (Hive, HBase, Pig, Flume, [Sqoop](#))  
Operational services which help manage the cluster ([Ambari](#), Falcon, and [Oozie](#))

Governance and integration	Data access	Security	Operations																																			
<p>Data workflow, lifecycle and governance</p> <p>Falcon <u>Sqoop</u> Flume NFS <u>WebHDFS</u></p>	<table><tr><td>Batch</td><td>Script</td><td>SQL</td><td>Nosql</td><td>Stream</td><td>Search</td><td>Others</td></tr><tr><td>Map-reduce</td><td>Pig</td><td>Hive/Tez, HCatalog</td><td>Database, <u>ApacheCassandra</u></td><td>Storm</td><td>Solr</td><td>Spark, in-memory, ISV engines</td></tr><tr><td colspan="7">YARN: data operating system</td></tr><tr><td colspan="7">HDFS (Hadoop Distributed File System)</td></tr><tr><td colspan="7">Data management</td></tr></table>	Batch	Script	SQL	Nosql	Stream	Search	Others	Map-reduce	Pig	Hive/Tez, HCatalog	Database, <u>ApacheCassandra</u>	Storm	Solr	Spark, in-memory, ISV engines	YARN: data operating system							HDFS (Hadoop Distributed File System)							Data management							<p>Authentication Authorization Accounting Data protection</p> <p>Storage: HDFS Resources: YARN Access: Hive, ... Pipeline: Falcon Cluster: Knox</p>	<p>Provision, manage, and monitor</p> <p><u>Ambari</u> <u>Zookeeper</u></p> <p>Scheduling</p> <p><u>Oozie</u></p>
Batch	Script	SQL	Nosql	Stream	Search	Others																																
Map-reduce	Pig	Hive/Tez, HCatalog	Database, <u>ApacheCassandra</u>	Storm	Solr	Spark, in-memory, ISV engines																																
YARN: data operating system																																						
HDFS (Hadoop Distributed File System)																																						
Data management																																						

# I Hadoop MapReduce

 Programming framework (library and runtime) for analysing datasets stored in HDFS

 Composed of user-supplied Map and Reduce functions:

- Map() - subdivide and conquer
- Reduce() - combine and reduce cardinality

Invoice	Date	Amount
1001	01-01-2016	\$100.00
1002	01-01-2016	\$95.00
1003	01-02-2016	\$100.00
1003	01-03-2016	\$75.00
1004	01-03-2016	\$50.00

**Map** Split data into Key/Value pairs

Key	Value
01-01-2016	{ \$100.00, \$95.00 }
01-02-2016	{ \$100.00 }
01-03-2016	{ \$75.00, \$50.00 }

Operate on values for each key

**Reduce**

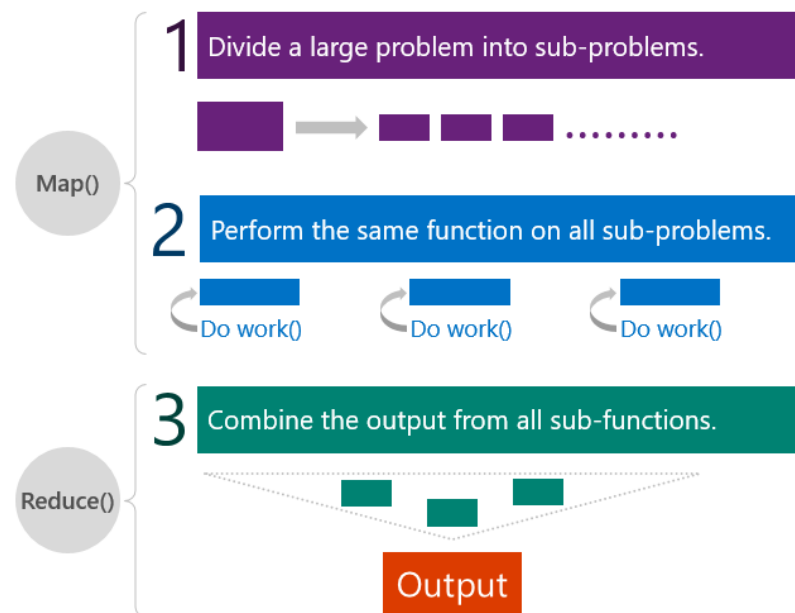
Key	Value
01-01-2016	$\Sigma = \$195.00$

Key	Value
01-02-2016	$\Sigma = \$100.00$

Key	Value
01-03-2016	$\Sigma = \$125.00$

**Output**


Key	Value
01-01-2016	\$195.00
01-02-2016	\$100.00
01-03-2016	\$125.00









# I Introduction to Big Data and HDInsight

## 1. What is HDInsight?

- a. Apache Hadoop on Azure
  - ✓ Hortonworks HDP on Azure VMs.
- b. Azure Storage or Azure Data Lake provides the HDFS layer.
- c. Azure SQL Database stores metadata.



## Data Processing and storage applications on HDP

Batch	Script	SQL	NoSQL	Stream	Other
Map/ reduce	Pig	Hive HCatalog	HBase	Storm	Mahout Oozie
C:/ 					

YARN resource manager

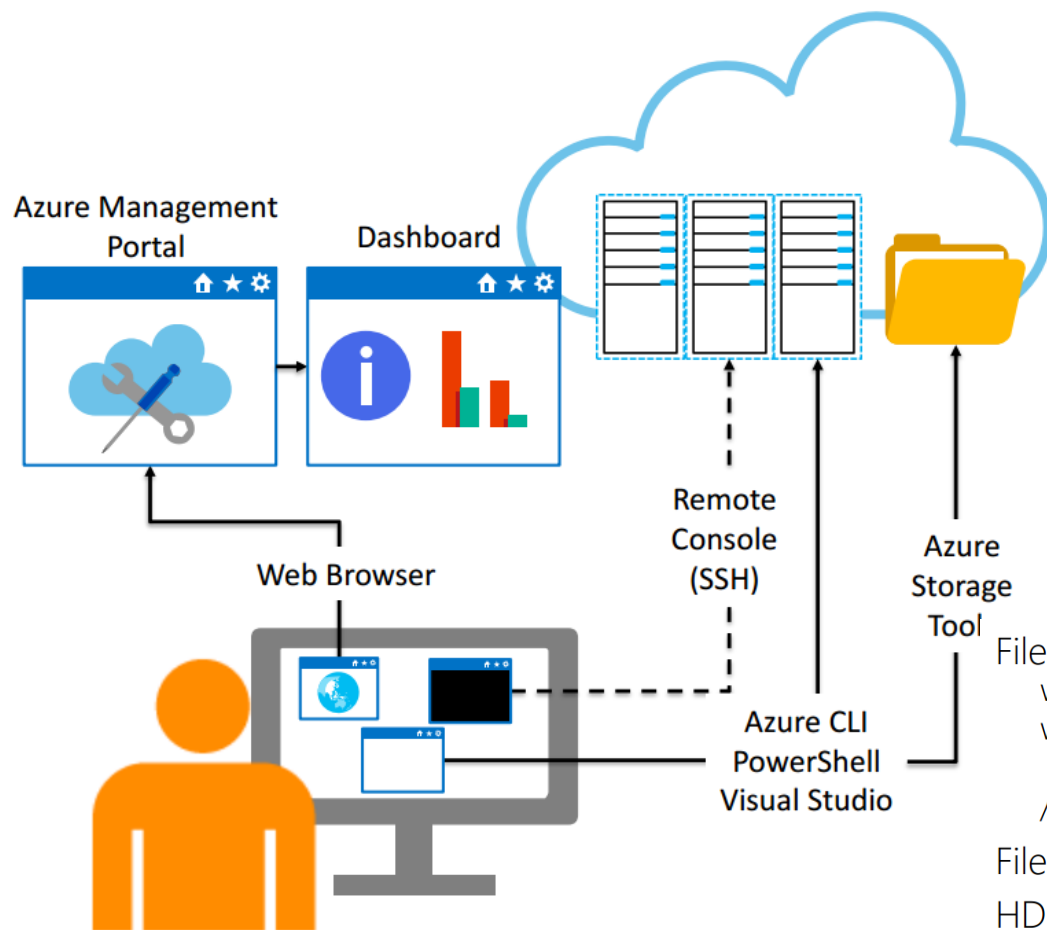
Hadoop distributed file system (HDFS)

- 🐘 Choose from a range of Hadoop-compatible tools
- 🐘 Run resource-intensive apps such as HBase or Storm on a separate cluster
- 🐘 Ensure better ROI by checking query results

**Demo:** Provisioning an HDInsight Cluster



# I HDInsight: What client tools can I use?



File paths can be referenced using WASB(S) or native syntax  
wasb://container@account.blob.core.windows.net/data/logs/file.txt  
wasb:///data/logs/file.txt (default storage account and container)  
or  
/data/logs/file.txt

File paths are case-sensitive

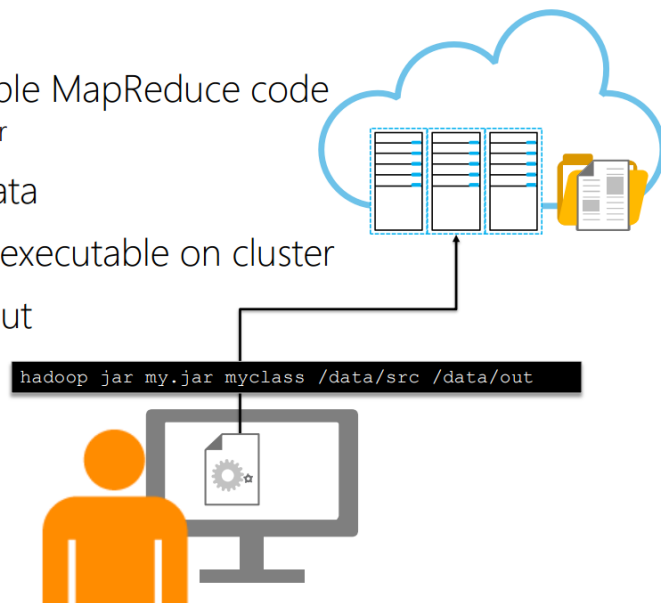
HDFS shell commands

- ls (list)
- cp and mv (copy and move)
- mkdir (make directory)
- rm and rm -r (remove and remove recursive)
- put and get (transfer files between local file system and HDFS)
- text, cat, and tail (display contents of file)



# I HDInsight: How do I Run a MapReduce Job

1. Compile executable MapReduce code  
Commonly a Java jar
2. Upload source data
3. Run MapReduce executable on cluster
4. Retrieve job output



```
Windows PowerShell ISE
File Edit View Tools Debug Add-ons Help

Provision Cluster.ps1 X
10 Login-AzureRmAccount
11
12 # Create a resource group
13 New-AzureRmResourceGroup -Name $resourceGro
14
15 # Create a storage account
16 Write-Host "Creating storage account..."
17 $storageAccountKey = Get-AzureRmStorageAcco
18 New-AzureRmStorageAccount -Name $storageAcc
19
20 # Create a Blob storage container
21 Write-Host "Creating container..."
22 $storageAccountKey = Get-AzureRmStorageAcco
23 $destContext = New-AzureStorageContext -Sto
24 New-AzureStorageContainer -Name $containerN

PS C:\Windows\system32>
```

The screenshot shows a Windows PowerShell ISE window. The main pane displays PowerShell commands for provisioning an Azure HDInsight cluster. The commands include logging in, creating a resource group, a storage account, and a blob storage container. The right pane shows the 'Commands' window with a list of modules and a search bar.

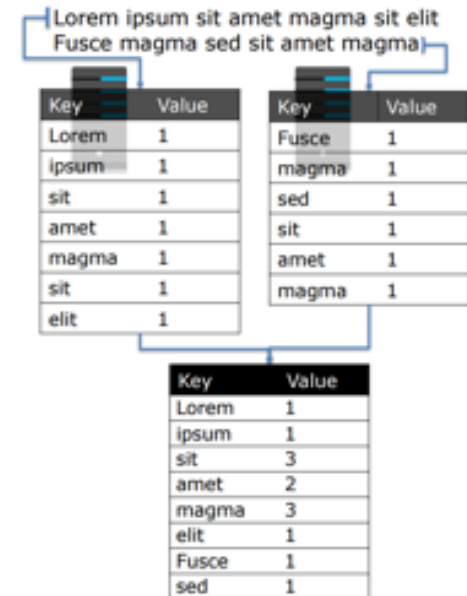

## Demo: Run a MapReduce Job from Windows

1. The Azure PowerShell module includes **cmdlets** to work with Azure services, including HDInsight
2. Use PowerShell to
  - a. Provision HDInsight Clusters
  - b. Upload/download files
  - c. Submit jobs
  - d. Manage cluster resources

# I Introduction to Big Data and HDInsight: Example

## 1. Word Count (The “Hello World” of MapReduce)

- Source text is divided among data nodes.
- Map phase generates key/value pairs with words as keys and placeholder values of 1.
- Reduce phase aggregates values for each key by adding the values for each word.



```
public static class Map extends Mapper<LongWritable, Text, Text, IntWritable> {  
    private final static IntWritable one = new IntWritable(1);  
    private Text word = new Text();  
    public void map(LongWritable key, Text value, Context context) {  
        String line = value.toString();  
        StringTokenizer tokenizer = new StringTokenizer(line);  
        while (tokenizer.hasMoreTokens()) {  
            word.set(tokenizer.nextToken());  
            context.write(word, one);  
        }  
    }  
}
```

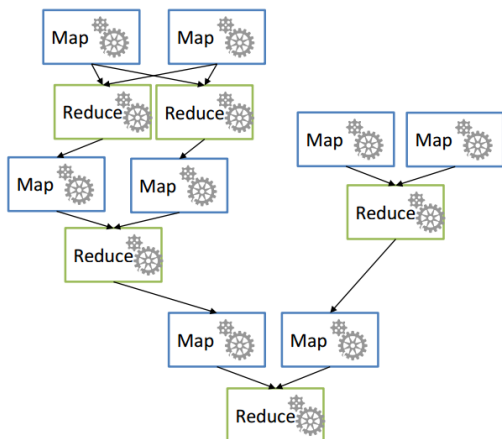
```
public static class Reduce extends Reducer<Text, IntWritable, Text, IntWritable> {  
    public void reduce(Text key, Iterable<IntWritable> values, Context context){  
        int sum = 0;  
        for (IntWritable val : values) {  
            sum += val.get();  
        }  
        context.write(key, new IntWritable(sum));  
    }  
}
```

# II Processing Big Data with Hive: What is Hive?


1. Apache HIVE is a data warehouse system for Hadoop.
2. A metadata service that **projects tabular schemas over folders**.
3. **Enables the contents of folders to be queried as tables**, using SQL-like query semantics, HiveQL.
4. Queries are translated into jobs
  - a. Execution engine can be **Tez** or **MapReduce**





```
set hive.execution.engine=mr;  
SELECT...
```



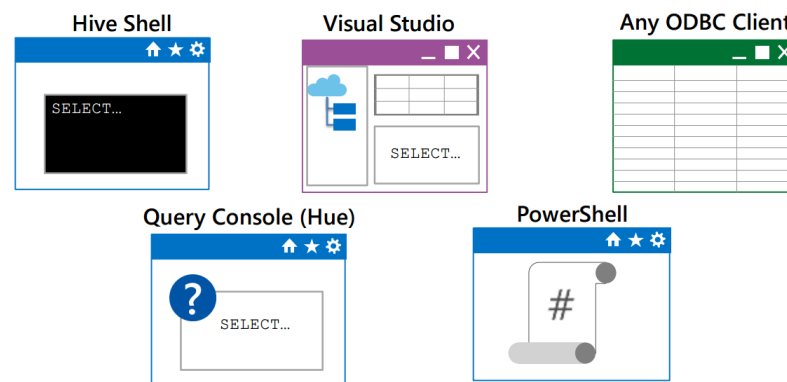
```
set hive.execution.engine=tez;  
SELECT...
```



**MapReduce vs. Tez**



Hive client tools include...



## II Hive: How do I create and load Hive tables?

1. Use the **CREATE TABLE** HiveQL statement
  - a. Defines schema metadata to be projected onto data in a folder when the table is queried (not when it is created)
2. Specify file **format** and **file location**
  - a. Defaults to textfile format in the <database>/<table-name> folder
    - ✓ Default database is in **/hive/warehouse**
    - ✓ Create additional databases using **CREATE DATABASE**
3. Create **internal** or **external** tables
  - a. Internal tables manage the lifetime of the underlying folders
  - b. External tables are managed independently from folders

```
CREATE TABLE table1
(col1 STRING,
 col2 INT)
ROW FORMAT DELIMITED FIELDS TERMINATED BY ' ';
```

Internal table (folders  
deleted when table is  
dropped)

Default location  
(/hive/warehouse/table1)

```
CREATE TABLE table2
(col1 STRING,
 col2 INT)
ROW FORMAT DELIMITED FIELDS TERMINATED BY ' '
STORED AS TEXTFILE LOCATION '/data/table2';
```

Stored in a custom folder (but  
still internal, so the folder is  
deleted when table is dropped)

```
CREATE EXTERNAL TABLE table3
(col1 STRING,
 col2 INT)
ROW FORMAT DELIMITED FIELDS TERMINATED BY ' '
STORED AS TEXTFILE LOCATION '/data/table3';
```

External table (folders and files  
are left intact in Azure Blob Store  
when the table is dropped)

## II Hive: Hive data types

---

1. Numeric
  - a. Integers: TINYINT, SMALLINT, INT, BIGINT
  - b. Fractional: FLOAT, DOUBLE, DECIMAL
2. Character
  - a. STRING, VARCHAR, CHAR
3. Date/Time
  - a. TIMESTAMP
  - b. DATE
4. Special
  - a. BOOLEAN, BINARY, ARRAY, MAP, STRUCT, UNIONTYPE

1. [Save data files in table folders](#) (or create table on existing files!)

```
PUT myfile.txt /data/table1
```

2. Use the [LOAD](#) statement

```
LOAD DATA (LOCAL) INPATH '/data/source' INTO TABLE MyTable;
```

3. User the [INSERT](#) statement

```
INSERT INTO TABLE Table2
```

```
SELECT Col1, UPPER(Col2)
```

```
FROM Table1;
```

4. Use a [CREATE TABLE AS SELECT](#) (CTAS) statement

```
CREATE TABLE Table3
```

```
ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t'
```

```
STORED AS TEXTFILE LOCATION '/data/summarytable'
```

```
AS
```

```
SELECT Col1, SUM(Col2) As Total
```

```
FROM Table1
```

```
GROUP BY Col1;
```

## II Hive: How do I query Hive Tables?

---

1. Query data using the **SELECT** HiveQL statement

```
SELECT Col1, SUM(Col2) AS TotalCol2
FROM MyTable
WHERE Col3 = 'ABC' AND Col4 < 10
GROUP BY Col1
ORDER BY Col4;
```

2. Hive translates the query into jobs and applies the table schema to the underlying data files
3. Views are named queries that abstract underlying tables

```
CREATE VIEW vSummarizeData
AS
SELECT col1, SUM(col2) AS TotalCol2
FROM mytable
GROUP BY col1;
```

```
SELECT col1, TotalCol2
FROM vSummarizeData
```


# II Hive: Partitioning, Skewing, and Clustering Tables

```
CREATE TABLE part_table
(col1 INT,
 col2 STRING)
PARTITIONED BY (col3 STRING);

INSERT INTO TABLE part_table PARTITION(col3='A')
SELECT col1, col2
FROM stg_table
WHERE col3 = 'A';


SET hive.exec.dynamic.partition = true;
SET hive.exec.dynamic.partition.mode=nonstrict;

INSERT INTO TABLE part_table PARTITION(col3)
SELECT col1, col2, col3
FROM stg_table;
```




```
CREATE TABLE skewed_table
(col1 INT,
 col2 STRING,
 col3 STRING)
SKEWED BY (col3) ON ('A') [STORED AS DIRECTORIES];

INSERT INTO TABLE skewed_table
SELECT col1, col2, col3
FROM stg_table;
```



```
CREATE TABLE clust_table
(col1 INT,
 col2 STRING,
 col3 STRING)
CLUSTERED BY (col3) INTO 3 BUCKETS;

INSERT INTO TABLE clust_table
SELECT col1, col2, col3
FROM stg_table;
```

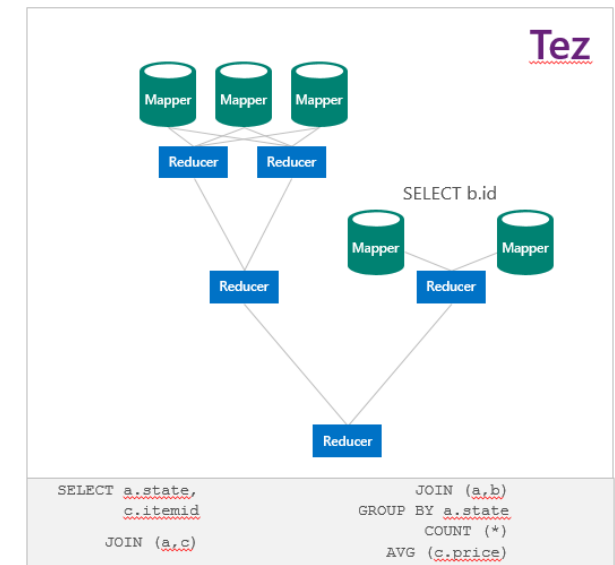
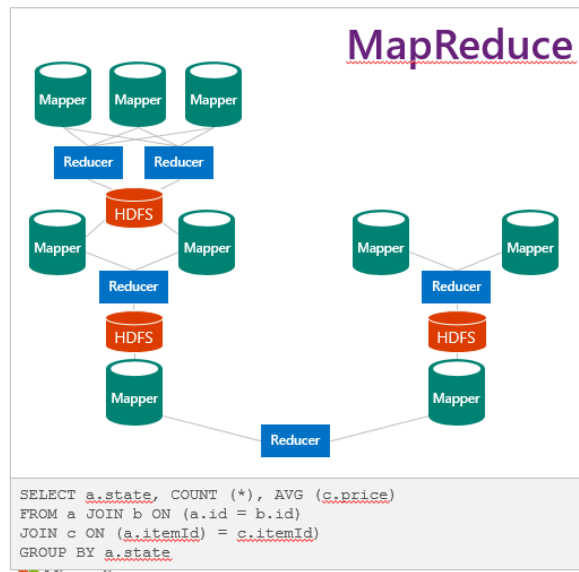




## II Hive Comparing RDBMS and Hive

	RDBMS	Hive
Structure	Schema On Write	Schema On Read
Access	SQL	SQL
Indexes	Yes	Yes
Updates	Yes	Yes (new in 0.14)
Locking	Yes	Table and Partition
Referential Integrity	Yes	No
Query Optimization	Yes	Yes

### Understanding Hive on Tez



# III Beyond Hive: What is Pig?

1. Pig is a high level scripting language used with Hadoop.
2. Pig performs **a series of transformations to data relations** based on Pig Latin statements.
3. Relations are loaded using **schema on read** semantics to project table structure at runtime.
4. You can run Pig Latin statements interactively in the **Grunt** shell or save a script file and run them as a batch.
5. A relation is an **outer bag**
  - a. A **bag** is a collection of tuples
  - b. A **tuple** is an ordered set of fields
  - c. A **field** is a data item
6. A **field** can contain an **inner bag**
7. A **bag** can contain **tuples** with nonmatching schema


{(a, 1)  
(b, 2)  
(c, 3)  
(d, {(4, 5), (6,7)})  
(e)  
(f, 8, 9)}

## Pig—Where it fits

Pig is designed to perform long series of data operations like:

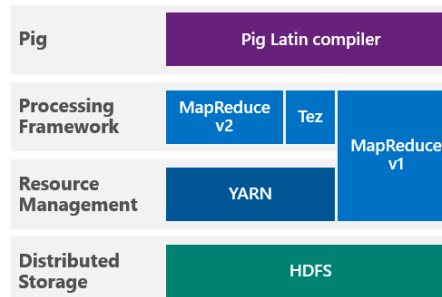
- Extract-transform-load (ETL) data pipelines
- Research on raw data
- Iterative data processing



### Key benefits

- Extensible.** Create custom functions to meet your particular processing requirements
- Easy to program.** Simplify and encode complex tasks involving interrelated data transformations as data flow sequences—easily write and maintain huge tasks
- Self-optimizing.** The system automatically optimizes execution of Pig jobs so you can focus on semantics

## What makes Pig so special?



Pigs are special because:

- Pigs eat anything.** Pig can operate on data whether it has metadata or not.
- Pigs live anywhere.** Pig is not tied to one particular parallel framework.
- Pigs are domestic animals.** Pig is designed to be easily controlled and modified by its users.
- Pigs fly.** Pig processes data quickly.

# III Pig: What kinds of things can I do with Pig?

Date, temperature

2013-06-01,12  
2013-06-01,14  
2013-06-01,16  
2013-06-02,9  
2013-06-02,12  
2013-06-02,9  
...



```
-- Load comma-delimited source data
Readings = LOAD '/weather/data.txt' USING PigStorage(',') AS (date:chararray, temp:long);
-- Group the tuples by date
GroupedReadings = GROUP Readings BY date;
-- Get the average temp value for each date grouping
GroupedAves = FOREACH GroupedReadings GENERATE group, AVG(Readings.temp) AS avgtemp;
-- Ungroup the dates with the average temp
AvgWeather = FOREACH GroupedAves GENERATE FLATTEN(group) as date, avgtemp;
-- Sort the results by date
SortedResults = ORDER AvgWeather BY date ASC;
-- Save the results in the /weather/summary folder
STORE SortedResults INTO '/weather/summary';
```



2013-06-01 14.00  
2013-06-02 10.00

## Common Pig Latin Operations


- LOAD
- FILTER
- FOR EACH ... GENERATE
- ORDER
- JOIN
- GROUP
- FLATTEN
- LIMIT
- DUMP
- STORE

# III Pig: What kinds of things can I do with Pig?

---

1. Pig generates Map and Reduce operations from Pig Latin
2. Jobs are generated on:
  - a. DUMP
  - b. STORE


```
Readings = LOAD '/weather/data.txt' USING PigStorage(',') AS (date, temp:long);
GroupedReadings = GROUP Readings BY date;
GroupedAvgs = FOREACH GroupedReadings GENERATE group, AVG(Readings.temp) AS avgtemp;
AvgWeather = FOREACH GroupedAvgs GENERATE FLATTEN(group) as date, avgtemp;
SortedResults = ORDER AvgWeather BY date ASC;
STORE SortedResults INTO '/weather/summary';
```



Job generated here



# III Pig: How do I run a Pig script?

1. Save a Pig Latin script file
2. Run the script using Pig  
`Pig wasb:///scripts/myscript.pig`
3. Consume the results using any Azure storage client
  - a. For example, Excel or Power BI
  - b. Default output does not include schema – just data



# III Pig: What are UDFs?

1. User Defined Functions (UDFs) extend the capability of Hive and Pig
2. Simpler than writing custom MapReduce components
3. Can be implemented using many languages, for example:
  - a. Java
  - b. C#
  - c. Python
4. Python is a (relatively) simple scripting language – ideal for UDFs
  - a. Intuitive syntax
  - b. Dynamic typing
  - c. Interpreted execution
5. Python is pre-installed on HDInsight clusters
  - a. Python 2.7 supports streaming from Hive
  - b. Jython (a Java implementation of Python) has native support in Pig



# III Pig: How do I use a Python UDF in Pig?

---

1. Pig natively supports Jython
  - a. Define the output schema as a Pig bag
  - b. Declare a Python function that receives an input parameter from Pig
  - c. Return results as fields based on the output schema

```
@outputSchema("result: {(a:chararray, b:int)}")
Def myfunction(i):
    ...

    return a, b
```

2. Use the Pig FOREACH...GENERATE statement to invoke a UDF

```
REGISTER 'wasb:///scripts/myscript.py' using jython as myscript;

src = LOAD '/data/source' AS (row:chararray);

res = FOREACH src GENERATE myscript.myfunction(row);
```

# III Pig: How do I use a Python UDF in Hive?

---

1. Hive exchanges data with Python using a [streaming](#) technique
  - a. Rows from Hive are passed to Python through [STDIN](#)
  - b. Processed rows from Python are passed to Hive through [STDOUT](#)

```
line = sys.stdin.readline()
```

```
...
```

```
print processed_row
```


2. Use the Hive [TRANSFORM](#) statement to invoke a UDF

```
add file wasb:///scripts/myscript.py;
```

```
SELECT TRANSFORM (col1, col2, col3)
  USING 'python myscript.py'
  AS(col1 string, col2 int, col3 string)
FROM mytable
ORDER BY col1;
```



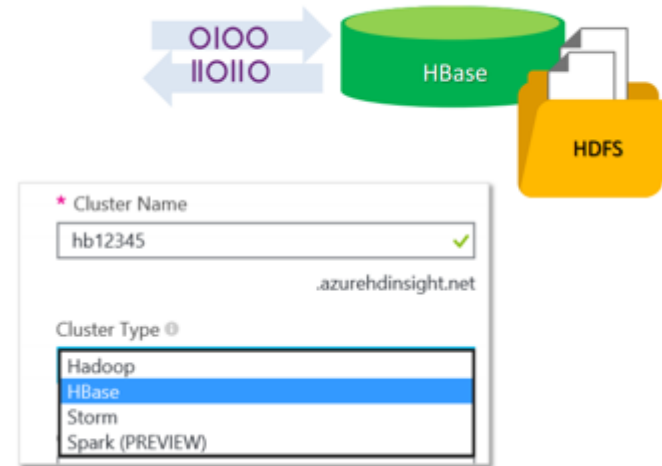
# III Hive and Pig – Complementing each other

	Hive	Pig
Language	SQL-like	PigLatin
Schemas/Types	Yes (explicit)	Yes (implicit)
Partitions	Yes	No
Server	Optional (Thrift)	No
User Defined Functions (UDF)	Yes (Java)	Yes (Java)
Custom Serializer/Deserializer	Yes	Yes
DFS Direct Access	Yes (implicit)	Yes (explicit)
Join/Order/Sort	Yes	Yes
Shell	Yes	Yes
Streaming	Yes	Yes
Web Interface	Yes	No
JDBC/ODBC	Yes (limited)	No
 Suitable Workload	Ad-Hoc Query	ETL

# IV HBase: What is Apache HBase?

1. A low-latency, NoSQL database built on Hadoop.
2. Modeled on Google's Big Table.
3. HBase stores data in StoreFiles on HDFS.

1. HDInsight supports an HBase cluster type
  - ✓ Choose Cluster Type in the Azure Portal.
2. Can be provisioned in a virtual network.



## What is HBase?

### Distributed, non-relational database

Columnar data model  
NoSQL on top of Hadoop

### Large scale

Linear scalability  
Billions of rows X millions of columns  
Many deployments with 1000+ nodes, PBs of data

### Low latency

Real-time random read/writes

### Open source

Modeled after Google's BigTable  
Started in 2006

### Key value store








Message systems  
Content management systems

### Examples

Facebook messages  
Twitter-like messages  
Webtable – web crawler/indexer

## HBase—low latency database

### Key benefits

-  Strongly consistent reads/writes
-  Automatic sharding
-  Automatic RegionServer failover
-  Hadoop/HDFS/MapReduce Integration
-  Java Client API
-  Supports Thrift and REST for non-Java front-ends
-  Block Cache and Bloom Filters—Operational management



# IV HBase: How Does HBase Store Data?

1. Data is stored as **key-value** pairs.
2. Table schema arranges values into column families.
3. Column family schema is flexible.
4. Columns are row-specific.

readings				
key	sensor		reading	
	id	location	datetime	value
1	Sensor1		2015-01-01	125.9
2	Sensor2		2015-01-01	152.3
3	Sensor1		2015-01-02	87.3
4	Sensor2		2015-01-02	151.8
5	Sensor1	Building 1	2015-01-03	126.3

1. Cells in a table are versioned.
2. Each versioned cell value is indicated by a timestamp.

readings				
key	sensor		reading	
	id	location	datetime	value
1	Sensor1		2015-01-01	125.9
2	Sensor2		2015-01-01	152.3
3	Sensor1		2015-01-02	87.3
4	Sensor2		2015-01-02	151.8
5	Sensor1	Building 1	2015-01-03	127.1
				147152436 126.3
				147152442 127.1

# IV HBase: How Do you work with an HBase Table?

```
create 'readings', 'sensor', 'reading'
```

readings				
key	sensor		reading	

```
put 'readings', '1', 'reading:value', '125.9'
```

readings				
key	sensor		reading	
	id		datetime	value
1	Sensor1		2015-01-01	125.9

```
put 'readings', '2', 'sensor:location', 'Building 2'
```

readings				
key	sensor		reading	
	id	location	datetime	value
1	Sensor1		2015-01-01	125.9
2	Sensor2	Building 2	2015-01-01	152.3

```
put 'readings', '2', 'reading:value', '157.6'
```

readings				
key	sensor		reading	
	id	location	datetime	value
1	Sensor1		2015-01-01	125.9
2	Sensor2	Building 2	2015-01-01	157.6

```
get 'readings', '2'
```

readings				
key	sensor		reading	
	id	location	datetime	value
1	Sensor1		2015-01-01	125.9
2	Sensor2	Building 2	2015-01-01	157.6

COLUMN	CELL
sensor:id	timestamp=142361, value=Sensor2
sensor:location	timestamp=142366, value=Building 2
reading:datetime	timestamp=142363, value=2015-01-01
reading:value	timestamp=142381, value=157.6

```
get 'readings', '2', {COLUMN => [reading:value]}
```

readings				
key	sensor		reading	
	id	location	datetime	value
1	Sensor1		2015-01-01	125.9
2	Sensor2	Building 2	2015-01-01	157.6

COLUMN	CELL
reading:value	timestamp=142379, value=152.3

# IV HBase: How Do you work with an HBase Table?

```
get 'readings', '2', {TIMERANGE => [0,142380]}
```

readings				
key	sensor		reading	
	id	location	datetime	value
1	Sensor1		2015-01-01	125.9
2	Sensor2	Building 2	2015-01-01	157.6

COLUMN	CELL
sensor:id	timestamp=142361, value=Sensor2
sensor:location	timestamp=142366, value=Building 2
reading:datetime	timestamp=142363, value=2015-01-01
reading:value	timestamp=142379, value=152.3

```
scan 'readings'
```

readings				
key	sensor		reading	
	id	location	datetime	value
1	Sensor1		2015-01-01	125.9
2	Sensor2	Building 2	2015-01-01	157.6

ROW	COLUMN+CELL
1	column=sensor:id, timestamp=142356, value=Sensor1
1	column=reading:datetime, timestamp=142357, value=2015-01-01
1	column=reading:value, timestamp=142359, value=125.9
2	column=sensor:id, timestamp=142361, value=Sensor2
2	column=sensor:location, timestamp=142366, value=Building 2
2	column=reading:datetime, timestamp=142363, value=2015-01-01
2	column=reading:value, timestamp=142381, value=157.6

# IV HBase: How Do you work with an HBase Table?

```
scan 'readings', {LIMIT => 1}
```

readings				
key	sensor		reading	
	id	location	datetime	value
1	Sensor1		2015-01-01	125.9
2	Sensor2	Building 2	2015-01-01	157.6
3	Sensor1	Building 1	2015-01-02	87.3

ROW	COLUMN+CELL
1	column=sensor:id, timestamp=142356, value=Sensor1
1	column=reading:datetime, timestamp=142357, value=2015-01-01
1	column=reading:value, timestamp=142359, value=125.9

```
scan 'readings', {STARTROW=>'2', STOPROW=>'3'}
```

readings				
key	sensor		reading	
	id	location	datetime	value
1	Sensor1		2015-01-01	125.9
2	Sensor2	Building 2	2015-01-01	157.6

ROW	COLUMN+CELL
2	column=sensor:id, timestamp=142361, value=Sensor2
2	column=sensor:location, timestamp=142366, value=Building 2
2	column=reading:datetime, timestamp=142363, value=2015-01-01
2	column=reading:value, timestamp=142375, value=157.6
3	column=sensor:id, timestamp=142371, value=Sensor1
3	column=sensor:location, timestamp=142372, value=Building 1
3	column=reading:datetime, timestamp=142373, value=2015-01-02

# IV HBase: How Do you work with an HBase Table?

```
delete 'readings', '2', 'sensor:location'
```

readings				
key	sensor		reading	
	id	location	datetime	value
1	Sensor1		2015-01-01	125.9
2	Sensor2		2015-01-01	157.6
3	Sensor1	Building 1	2015-	
4	Sensor2	Building 2	2015-	
5	Sensor1	Building 1	2015-	
6	...			

```
drop 'readings'
```

```
deleteall 'readings', '4'
```


readings				
key	sensor			
	id	location		
1	Sensor1			
2	Sensor2		2015-01-01	157.6
3	Sensor1	Building 1	2015-01-02	87.3
5	Sensor1	Building 1	2015-01-03	126.3
6	...			

readings				
key	sensor		reading	
	id	location	datetime	value
1	Sensor1		2015-01-01	125.9
2	Sensor2		2015-01-01	157.6
3	Sensor1	Building 1	2015-01-02	87.3
5	Sensor1	Building 1	2015-01-03	126.3
6	...			


# IV HBase: How to use HBase with Hive?

## How Do you Bulk Load Data into HBase?

1. Upload data to HDFS
  - ✓ In Azure Storage
2. Import into a StoreFile
3. Load the StoreFile to an HBase table



## How Do You Query HBase Tables from Hive?



Load HBase from HDFS  
Load HBase from HBase

```
CREATE EXTERNAL TABLE hivetable
(key STRING, col1 STRING, col2 STRING)
STORED BY 'org.apache.hadoop.hive.hbase.HBaseStoragehandler'
WITH SERDEPROPERTIES
('hbase.columns.mapping' = ':key,cf:col1, cf:col2')
TBLPROPERTIES('hbase.table.name' = 'hhtable')
```



# IV Sqoop: What is Sqoop?

1. Sqoop is a database integration service
  - a. Built on open source Hadoop technology
  - b. Enables bi-directional data transfer between Hadoop clusters and databases via JDBC



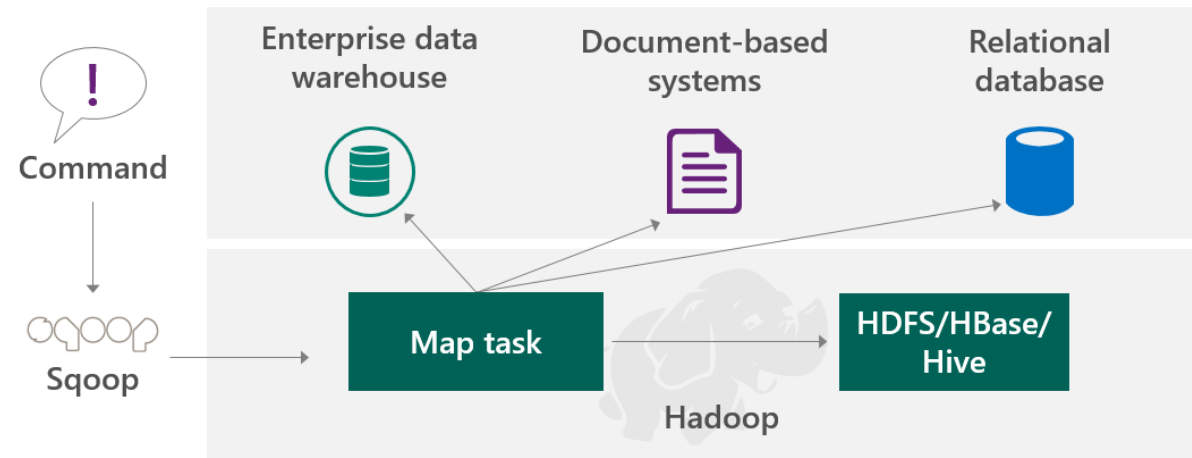
2. Basic syntax

`sqoop command --arg1, --arg2, ..., --argN`

3. Commands

- a. import / export
- b. help
- c. import-all-tables
- d. create-hive-tables
- e. list-databases / list-tables
- f. eval
- g. codegen
- h. version

Sqoop is designed to efficiently transfer bulk data between Apache Hadoop and structured datastores such as relational databases



## IV Sqoop: How do I run Sqoop commands?

---

### 1. **sqoop import**

**--connect jdbc-connection-string**  
**--username user-name --password password | -P**  
**--table table-name --columns col,...,colN | --query 'SELECT...'**  
**--warehouse-dir | --target-dir path**  
**--fields-terminated-by char --lines-terminated-by char**  
**--hive-import [--hive-overwrite]**  
**-m | --num-mappers number-of-mappers**


### 2. **sqoop export**

**--connect jdbc-connection-string**  
**--username user-name --password password | -P**  
**--table table-name**  
**--export-dir path**  
**--fields-terminated-by char --lines-terminated-by char**  
**-m | --num-mappers number-of-mappers**

# IV: Oozie and Mahout

## Oozie workflow

- Consists of workflow.xml file and the necessary files for the workflow itself
- Workflow saved to WASB store
- Control flow nodes for determining the execution path
- Action nodes for execution the job or tasks



## Mahout

Mahout is an Apache project to implement machine learning algorithms in Hadoop

Uses Hadoop to distribute algorithms to parallelize large scale workloads



### Three different categories of algorithms

- Recommendation:** Uses user information and community information to build recommendation. E.g. Netflix, Amazon, Pandora
- Classification:** Uses known data to classify new data into known buckets. E.g. Spam Detection
- Clustering:** Forms groups of data into similar categories. E.g. Bing News

